

CS 3305B

Brief Intro to Threads

Lecture 6

Jan 25 2017

Introduction

- ❑ Multiple applications run concurrently!
- ❑ This means that there are multiple processes running on a computer

A "Single Threaded" Program

```
main()
```

```
{
```


```
.....
```

```
{
```

Introduction

- ❑ Applications often need to perform many tasks at once
- ❑ This requires multiple **threads** of execution within a single process

Example




□ Example: Word processor

□ Tasks include:

- Display graphics
- Respond to keystrokes from the user
- Perform spelling and grammar checking

Example

- ❑ Example: Web server
 - ❑ It is desirable to service requests concurrently




Introduction

- ❑ Earlier we discussed the use of forking to create a process
- ❑ For example we could
 - ❑ **Word processor example:** fork a process for each task
 - ❑ **Web server example:** fork a process for each request
- ❑ Not very efficient since a fork copies everything

Threads


- ❑ A **thread** is a basic unit of CPU utilization
- ❑ Threads of a process share memory but can execute independently
- ❑ A traditional process can be viewed as a memory address space with a single thread

Thread Usage - Word Processor



A word processor program with three threads.

Thread Usage - Web Server



Why Not Fork?

- ❑ You certainly can fork a new process
- ❑ In fact, the first implementation of Apache web servers (Apache 1.0) forked N processes when the web server was started
 - ❑ "N" was defined in a configuration file
 - ❑ Each child process handled one connection at a time
- ❑ **Problem:** Process creation is time consuming and resource intensive
- ❑ Creating threads is not as expensive

Why Not Fork?

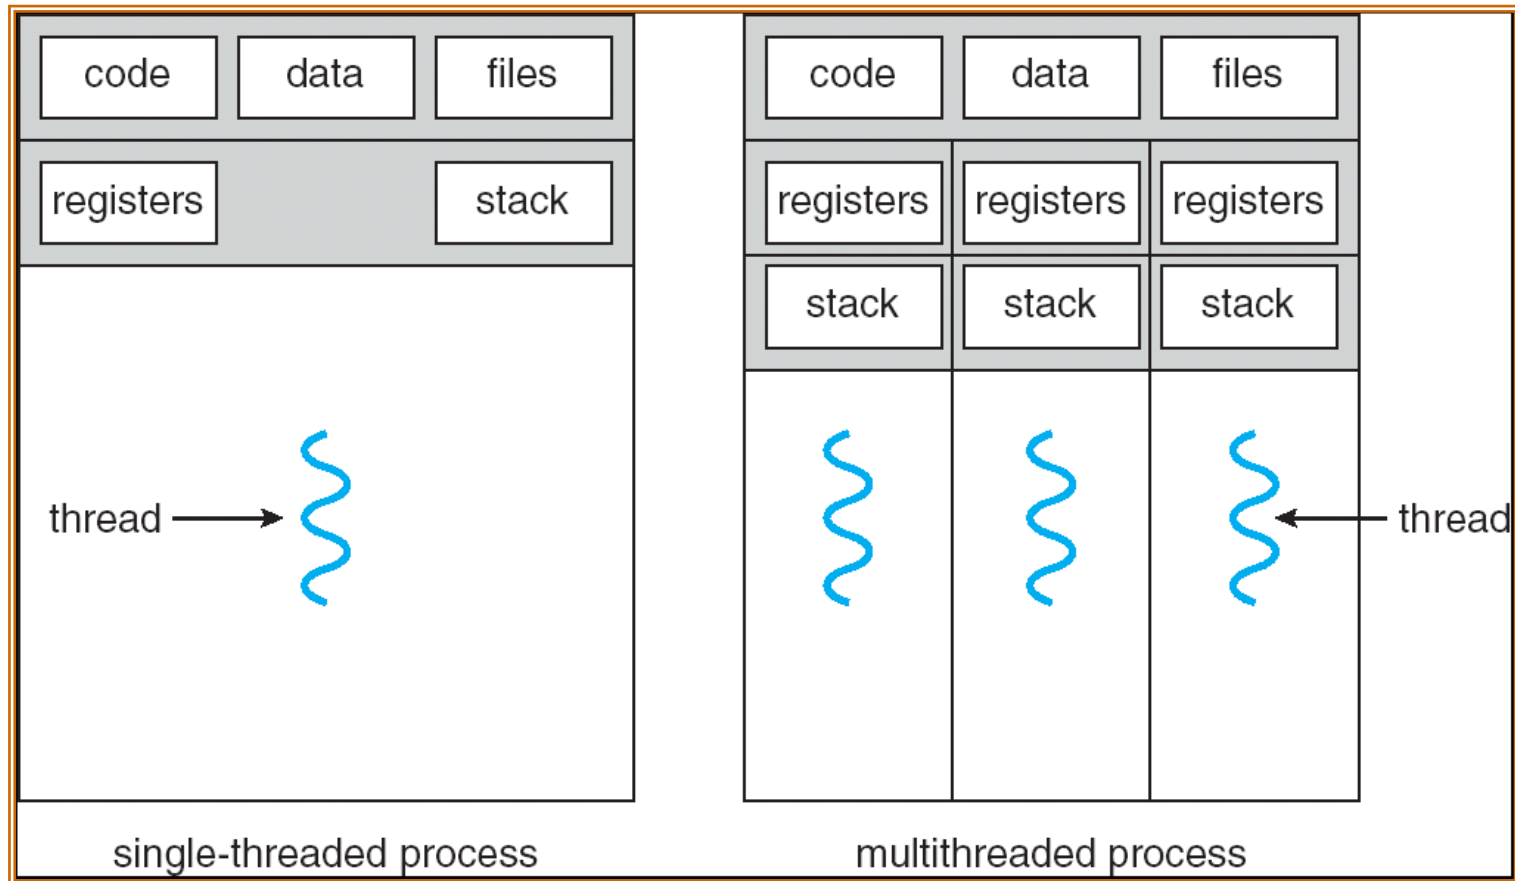
- ❑ Let's look at web servers
 - ❑ This allowed a child process to handle multiple connections at a time
 - ❑ Web servers have caches for read pages
 - ❑ Forking means that these caches cannot be shared
 - ❑ Using threads allows for these caches to be shared

Thread State

- ❑ Threads share
 - ❑ Process address space
 - ❑ Text
 - ❑ Data (global variables)
 - ❑ Heap (dynamically allocated data)
 - ❑ OS state
 - ❑ Open files, sockets, locks
- ❑ Threads have their own CPU context
 - ❑ Program counter(PC), Stack pointer (SP), register state, stack

Single and Multithreaded Processes

they share it



Benefits of Threads

- ❑ Responsiveness

they have their own cpu context

- ❑ Overlap computation and blocking due to I/O on a single CPU

- ❑ Resource Sharing

- ❑ Example: Word processor

- ❑ The document is shared in memory.
 - ❑ Forking would require replication

- ❑ Allocating memory and resources for process creation is costly

- ❑ Context-switching is faster

Thread Libraries

- ❑ A thread library provides the programmer with an *API* for creating and managing threads
- ❑ Three main libraries:
 - ❑ POSIX Pthreads
 - ❑ Win32
 - ❑ Java

Problem

- ❑ Sharing global variables is dangerous - two threads may attempt to modify the same variable at the same time.
- ❑ Use support for mutual exclusion primitives that can be used to protect against this problem.
- ❑ The general idea is to **lock** something before accessing global variables and to **unlock** as soon as you are done.
- ❑ More on this topic later in the course

Summary

- ❑ Introduction to the concept of threads
- ❑ There will be more discussion